

Graduate Descent

(<http://timvieira.github.io/blog/>)

- Atom (</blog/atom.xml>)
- About (<http://timvieira.github.io/>)
- Archive (</blog/archives.html>)
- Misc (<http://timvieira.github.io/blog/category/misc.html>)

Exp-normalize trick

Feb 11, 2014

This trick is the very close cousin of the infamous log-sum-exp trick (`scipy.misc.logsumexp` (<http://docs.scipy.org/doc/scipy/reference/generated/scipy.misc.logsumexp.html>)),
Supposed you'd like to evaluate a probability distribution π parametrized by a vector $\mathbf{x} \in \mathbb{R}^n$ as follows:

$$\pi_i = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)}$$

The exp-normalize trick leverages the following identity to avoid numerical overflow. For any $b \in \mathbb{R}$,

$$\pi_i = \frac{\exp(x_i - b) \exp(b)}{\sum_{j=1}^n \exp(x_j - b) \exp(b)} = \frac{\exp(x_i - b)}{\sum_{j=1}^n \exp(x_j - b)}$$

In other words, the π is shift-invariant. A reasonable choice is $b = \max_{i=1}^n x_i$. With this choice, overflow due to exp is impossible—the largest number exponentiated after shifting is 0.

Exp-normalize v. log-sum-exp

If what you want to remain in log-space, that is, compute $\log(\pi)$, you should use `logsumexp`. However, if π is your goal, then exp-normalize trick is for you! Since it avoids additional calls to exp, which would be required if using log-sum-exp and more importantly exp-normalize is more numerically stable!

Log-sum-exp for computing the log-distribution

$$\log \pi_i = x_i - \text{logsumexp}(\mathbf{x})$$

where

$$\text{logsumexp}(\mathbf{x}) = b + \log \sum_{j=1}^n \exp(x_j - b)$$

Typically with the same choice for b as above.

Numerically-stable sigmoid function

The sigmoid function can be computed with the exp-normalize trick in order to avoid numerical overflow. In the case of $\text{sigmoid}(x)$, we have a distribution with unnormalized log probabilities $[x, 0]$, where we are only interested in the probability of the first event. From the exp-normalize identity, we know that the distributions $[x, 0]$ and $[0, -x]$ are equivalent (to see why, plug in $b = \max(0, x)$). This is why sigmoid is often expressed in one of two equivalent ways:

$$\text{sigmoid}(x) = 1/(1 + \exp(-x)) = \exp(x)/(\exp(x) + 1)$$

Interestingly, each version covers an extreme case: $x = \infty$ and $x = -\infty$, respectively. Below is some python code which implements the trick:

```
def sigmoid(x):
    "Numerically-stable sigmoid function."
    if x >= 0:
        z = exp(-x)
        return 1 / (1 + z)
    else:
        # if x is less than zero then z will be small, denom can't be
        # zero because it's 1+z.
        z = exp(x)
        return z / (1 + z)
```



Comments

0 Comments

Graduate Descent

 Login ▾

 Recommend 1

 Share

Sort by Best ▾



Start the discussion...

Be the first to comment.

ALSO ON GRADUATE DESCENT

Dimensional analysis of gradient ascent

2 comments • 10 months ago •



Tim Vieira — Adaptive step size methods, e.g. Adagrad, Adadelta, and Adam, are generic optimization methods. Thus, they can be used to

Complex-step derivative

5 comments • 3 years ago •



Tim Vieira — Thanks "other Tim". I've heard great things about Kmett's work on AD from my office mate. The Conal article looks awesome -- it

KL-divergence as an objective function

3 comments • 2 years ago •






Rasmus Arnling Bååth — Just what I needed! Thanks!

Conditional random fields as Deep learning models?

4 comments • 2 years ago •




Tim Vieira — Thanks for the pointer!

 Subscribe  Add Disqus to your site Add Disqus Add  Privacy

Recent Posts

- Counterfactual reasoning and learning from logged data (<http://timvieira.github.io/blog/post/2016/12/19/counterfactual-reasoning-and-learning-from-logged-data/>)
- Heaps for incremental computation (<http://timvieira.github.io/blog/post/2016/11/21/heaps-for-incremental-computation/>)
- Reversing a sequence with sublinear space (<http://timvieira.github.io/blog/post/2016/10/01/reversing-a-sequence-with-sublinear-space/>)
- Evaluating $\nabla f(x)$ is as fast as $f(x)$ (<http://timvieira.github.io/blog/post/2016/09/25/evaluating-fx-is-as-fast-as-fx/>)
- Fast sigmoid sampling (<http://timvieira.github.io/blog/post/2016/07/04/fast-sigmoid-sampling/>)

Tags

datastructures (<http://timvieira.github.io/blog/tag/datastructures.html>), calculus (<http://timvieira.github.io/blog/tag/calculus.html>), statistics (<http://timvieira.github.io/blog/tag/statistics.html>), importance-sampling (<http://timvieira.github.io/blog/tag/importance-sampling.html>), structured-prediction (<http://timvieira.github.io/blog/tag/structured-prediction.html>), crf (<http://timvieira.github.io/blog/tag/crf.html>), decision-making (<http://timvieira.github.io/blog/tag/decision-making.html>), misc (<http://timvieira.github.io/blog/tag/misc.html>), sampling 

(<http://timvieira.github.io/blog/tag/sampling.html>), deep-learning (<http://timvieira.github.io/blog/tag/deep-learning.html>), rant (<http://timvieira.github.io/blog/tag/rant.html>), optimization (<http://timvieira.github.io/blog/tag/optimization.html>), algorithms (<http://timvieira.github.io/blog/tag/algorithms.html>), visualization (<http://timvieira.github.io/blog/tag/visualization.html>), machine-learning (<http://timvieira.github.io/blog/tag/machine-learning.html>), Gumbel (<http://timvieira.github.io/blog/tag/gumbel.html>), autodiff (<http://timvieira.github.io/blog/tag/autodiff.html>), numerical (<http://timvieira.github.io/blog/tag/numerical.html>), counterfactual-reasoning (<http://timvieira.github.io/blog/tag/counterfactual-reasoning.html>)

GitHub Repos

- [timvieira.github.com](https://github.com/timvieira/timvieira.github.com) (<https://github.com/timvieira/timvieira.github.com>)
Personal webpage.
- [arsenal](https://github.com/timvieira/arsenal) (<https://github.com/timvieira/arsenal>)
Arsenal of python utilities.
- [skid](https://github.com/timvieira/skid) (<https://github.com/timvieira/skid>)
bookmarks, simply kept in directories.
- [learning-to-prune](https://github.com/timvieira/learning-to-prune) (<https://github.com/timvieira/learning-to-prune>)
Learning to Prune: Exploring the Frontier of Fast and Accurate Parsing

@timvieira (<https://github.com/timvieira>) on GitHub

Latest Tweets

- Status updating...

Follow @xtimv { 616 followers }

Copyright © 2014-2016 - Tim Vieira - Powered by Pelican (<http://getpelican.com>)

