

Lecture 2:

Caffe: getting started

Forward propagation


boris.ginzburg@intel.com

Agenda

- Caffe – getting started
- Test description
- Network topology definition
- Basic layers: definition and forward propagation
 - Convolutional
 - Pooling
 - ReLU
 - Fully Connected layer
 - Softmax
- Implementation details of Convolutional layer
- MNIST training

Open-source Deep Learning libraries

1. <http://caffe.berkeleyvision.org/>
Very fast. C++/ CUDA, Python and Matlab wrappers
2. <https://code.google.com/p/cuda-convnet2/>
Just released. Excellent tutorial. Best Cuda code.
3. <http://torch.ch/>
Excellent tutorial, C++/Cuda, Lua.
4. <http://deeplearning.net/software/pylearn2/>:
Integrated with Theano, C++/Cuda, Python
5. <http://torontodeeplearning.github.io/convnet/>
C++/CUDA.



Caffe:

Convolutional Architecture for Fast Feature Embedding

Created by Yangqing Jia

Developed by BVLC

caffe.berkeleyvision.org

bvlc.eecs.berkeley.edu





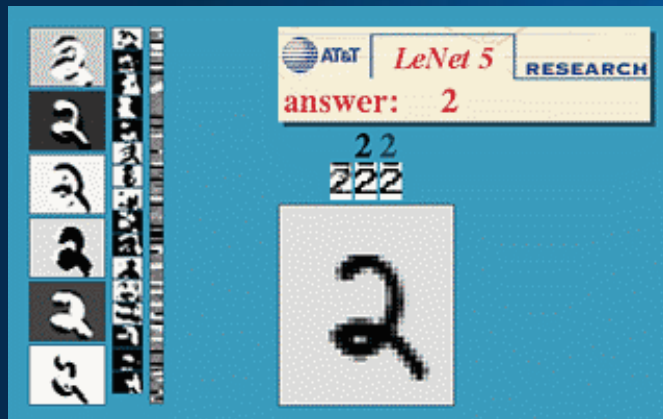
Caffe: installation

1. Ubuntu 12.04
2. Cuda 5.5 or 6.0
 - SDK - required, NVidia card is optional ☺
3. BLAS:
 - OpenBLAS or Intel MKL(Math Kernel Lib)

\$ git clone <https://github.com/BVLC/caffe>

Caffe: example 1 - MNIST

- Database: <http://yann.lecun.com/exdb/mnist/>
- Demo: <http://yann.lecun.com/exdb/lenet/index.html>



Caffe: database format

src/tools/convert_mnist_data.cpp: MNIST format → leveldb

1. leveldb: <https://code.google.com/p/leveldb/>

- <key,value>: arbitrary byte arrays; data is stored sorted by key; callers can provide a custom comparison function to override the sort order.
- basic operations : Put(key,value), Get(key), Delete(key).

2. caffe “dev” branch supports lmdb: <http://symas.com/mdb/>

- <key;value> ; data is stored sorted by key
- uses **memory-mapped files**: the read performance of a pure in-memory db while still offering the persistence of standard disk-based db
- concurrent

Caffe: configuration files

1. Solver descriptor:

▶ http://caffe.berkeleyvision.org/mnist_solver_prototxt.html

1. Net descriptor:

▶ http://caffe.berkeleyvision.org/mnist_prototxt.html

Parameters are defined in src/caffe/proto/caffe.proto.

Protobuf (Google protocol buffers) format - easy-to-use automatic generation of configuration files:

<https://developers.google.com/protocol-buffers/docs/overview>

LeNet Topology


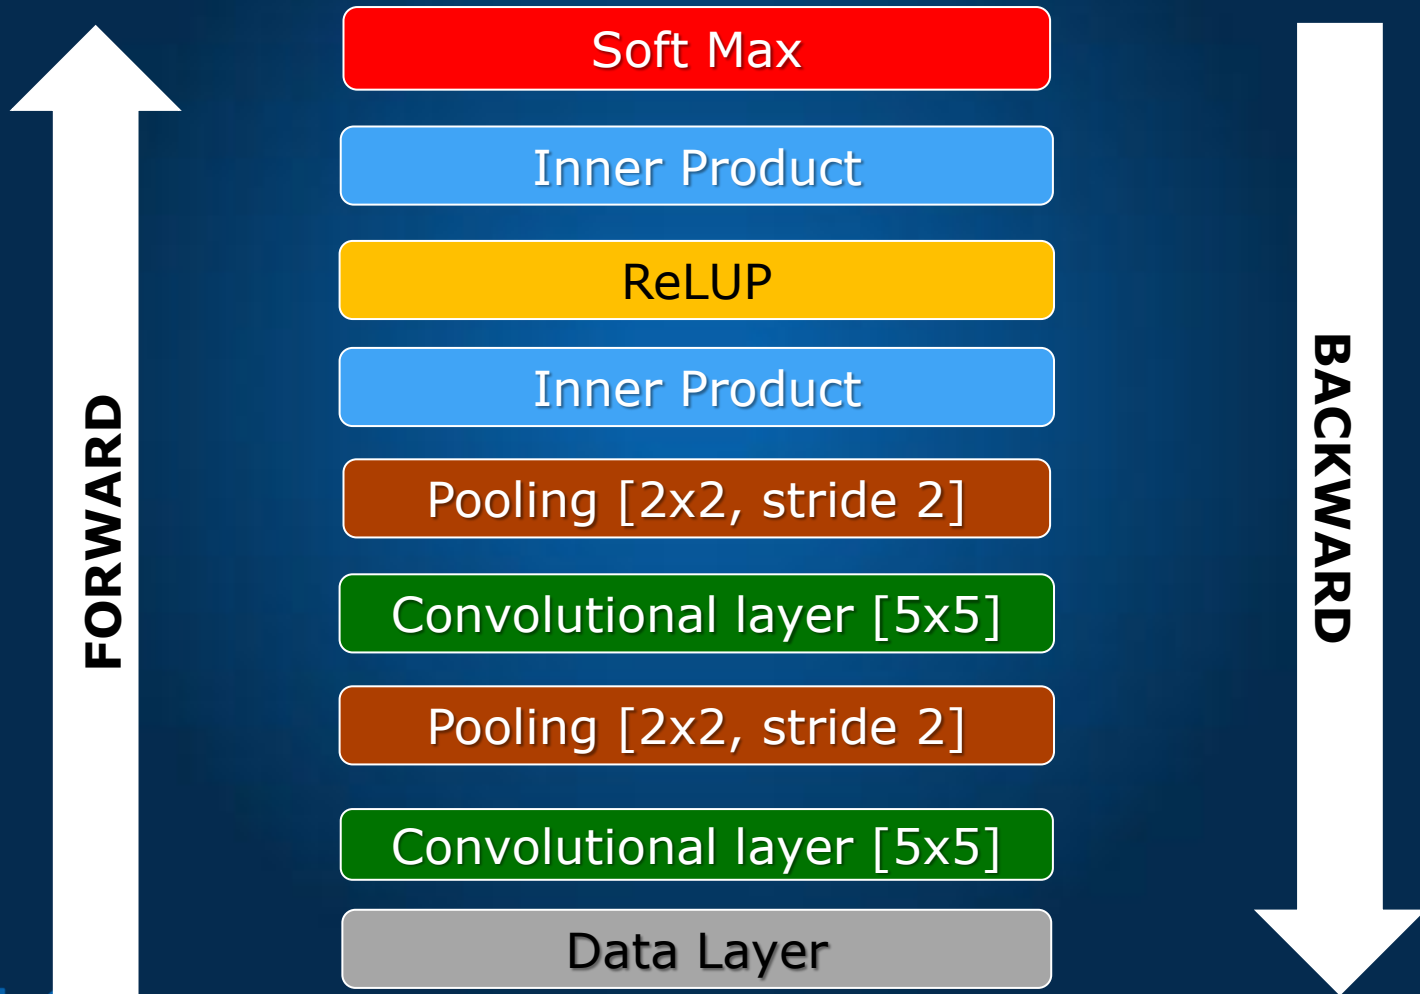


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>

LeNet topology



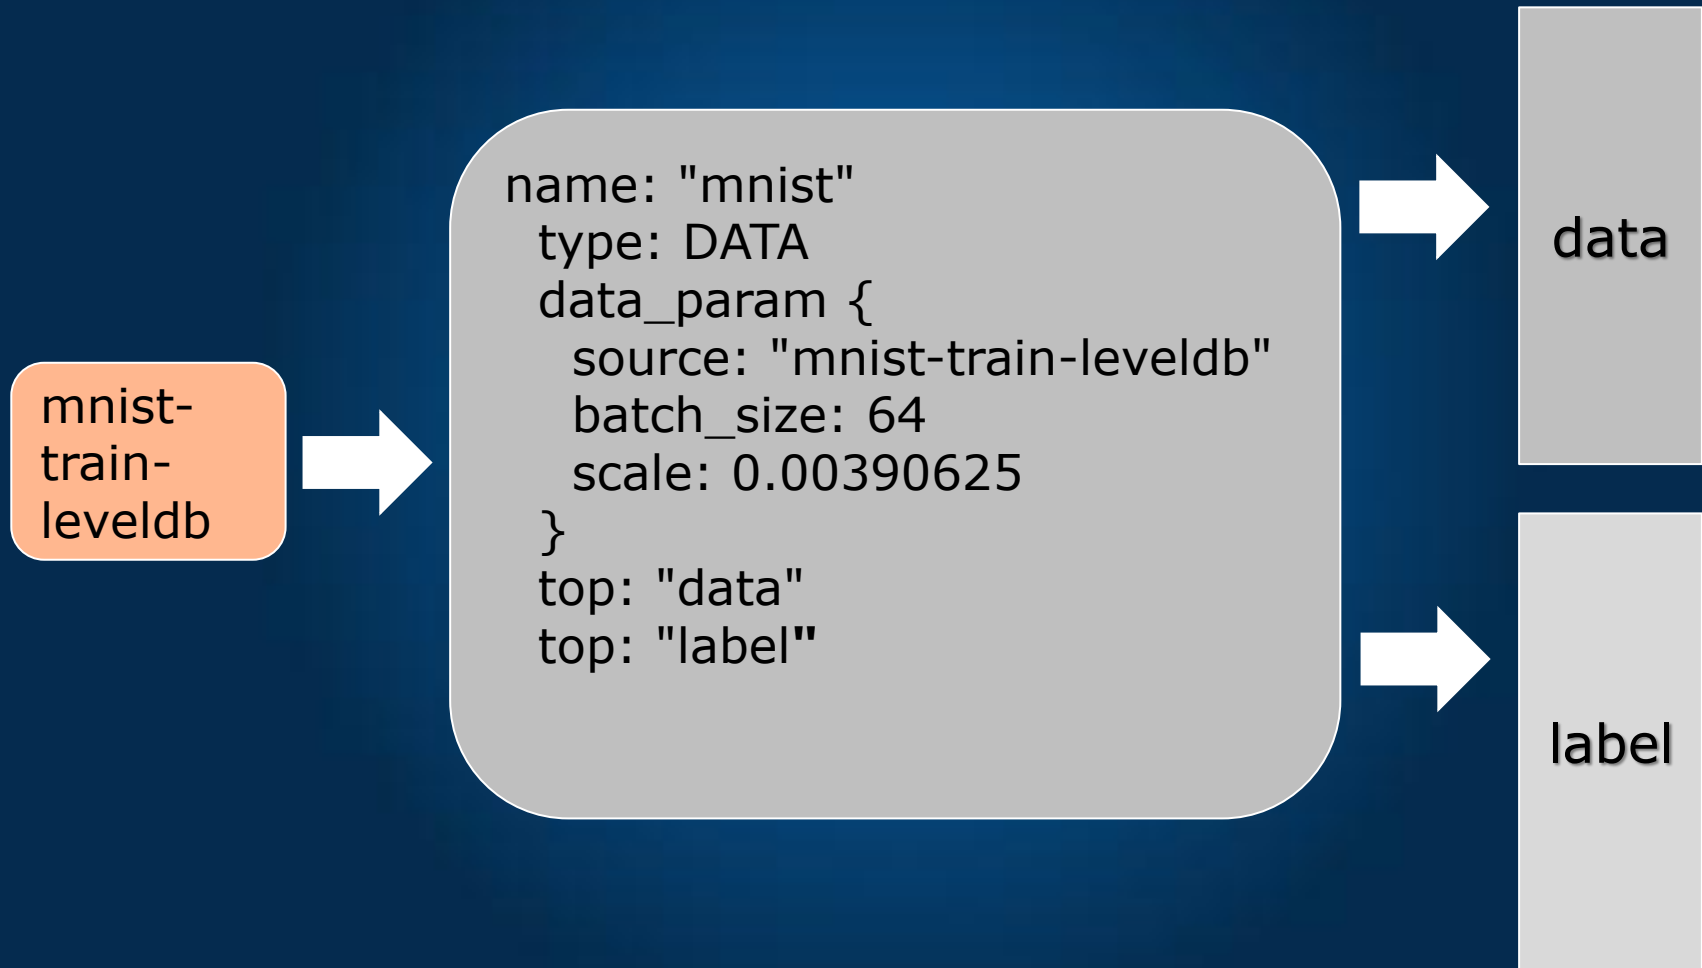
Layer:: Forward()

```
class Layer {  
    Setup (bottom, top);    // initialize layer  
    Forward (bottom, top);  //compute next layer  
    Backward( top, bottom); //compute gradient  
}
```

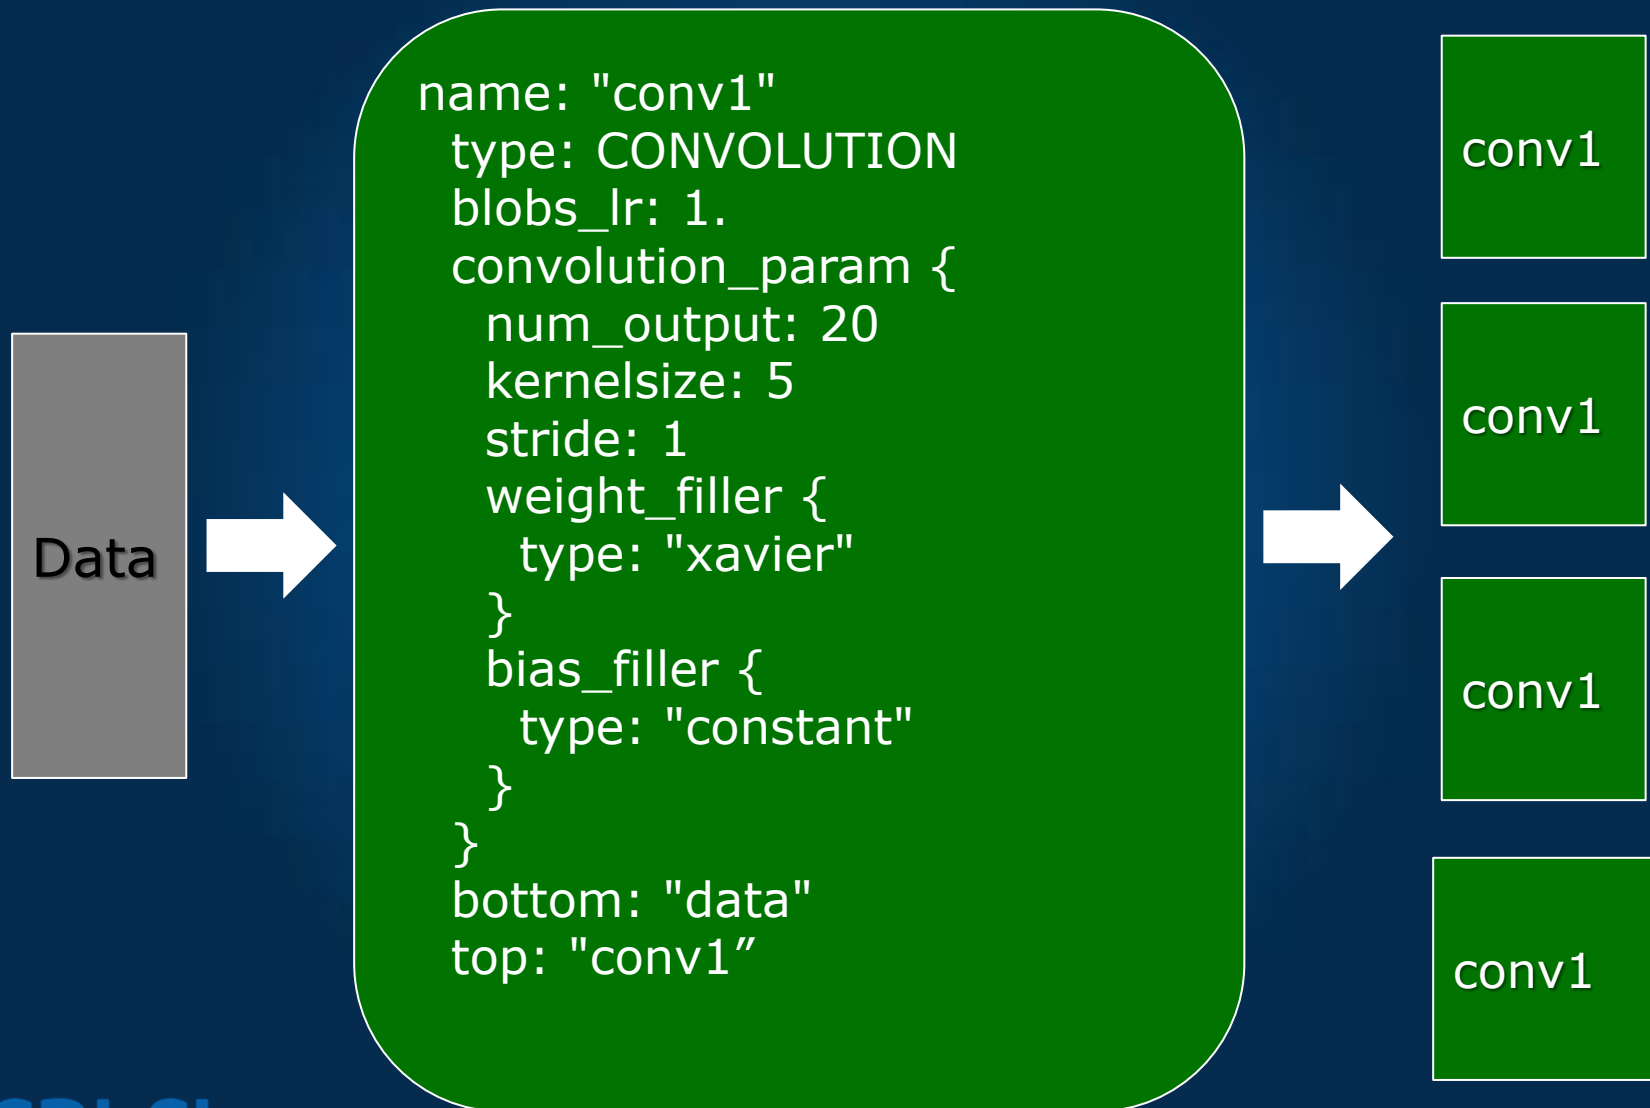
Layer::Forward() propagate y_{l-1} to next layer:

$$y_l = f(w_l, y_{l-1})$$

Data Layer




Convolutional Layer



Convolutional Layer


```
for (n = 0; n < N; n++)  
  for (m = 0; m < M; m++)  
    for (y = 0; y < Y; y++)  
      for (x = 0; x < X; x++)  
        for (p = 0; p < K; p++)  
          for (q = 0; q < K; q++)  
             $y_L(n; x, y) += y_{L-1}(m, x+p, y+q) * w(m, n; p, q);$ 
```

Add bias...



Pooling Layer

```
name: "pool1"  
type: POOLING  
pooling_param {  
  kernel_size: 2  
  stride: 2  
  pool: MAX  
}  
bottom: "conv1"  
top: "pool1"
```



for (p = 0; p < k; p++)

for (q = 0; q < k; q++)

$$y_L(x, y) = \max (y_L(x, y), y_{L-1}(x*s + p, y*s + q));$$

Pooling helps to extract features that are increasingly invariant to local transformations of the input image.

Inner product (Fully Connected) Layer

```
name: "ip1"  
type: INNER_PRODUCT  
blobs_lr: 1.  
blobs_lr: 2.  
inner_product_param {  
  num_output: 500  
  weight_filler {  
    type: "xavier"  
  }  
  bias_filler {  
    type: "constant"  
  }  
}  
bottom: "pool2"  
top: "ip1"
```

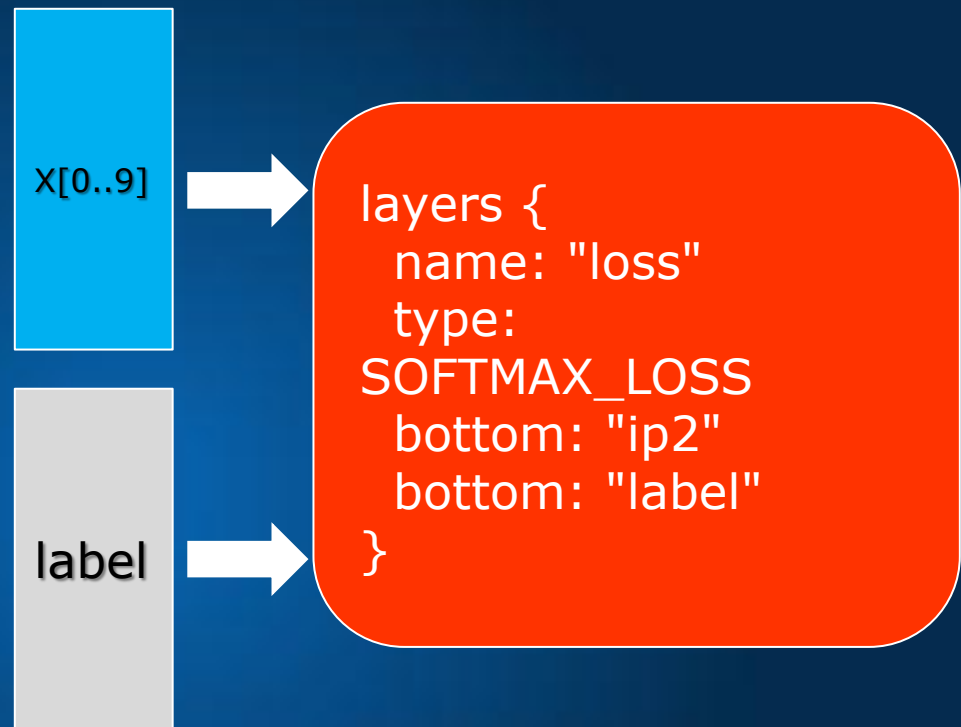
$$Y_L(n) = \sum W_L(n, m) * Y_{L-1}(m)$$

ReLU Layer

```
layers {  
  name: "relu1"  
  type: RELU  
  bottom: "ip1"  
  top: "ip1"  
}
```

$$Y_L(n; x, y) = \max(Y_{L-1}(n; x, y), 0);$$

SoftMax + Loss Layer



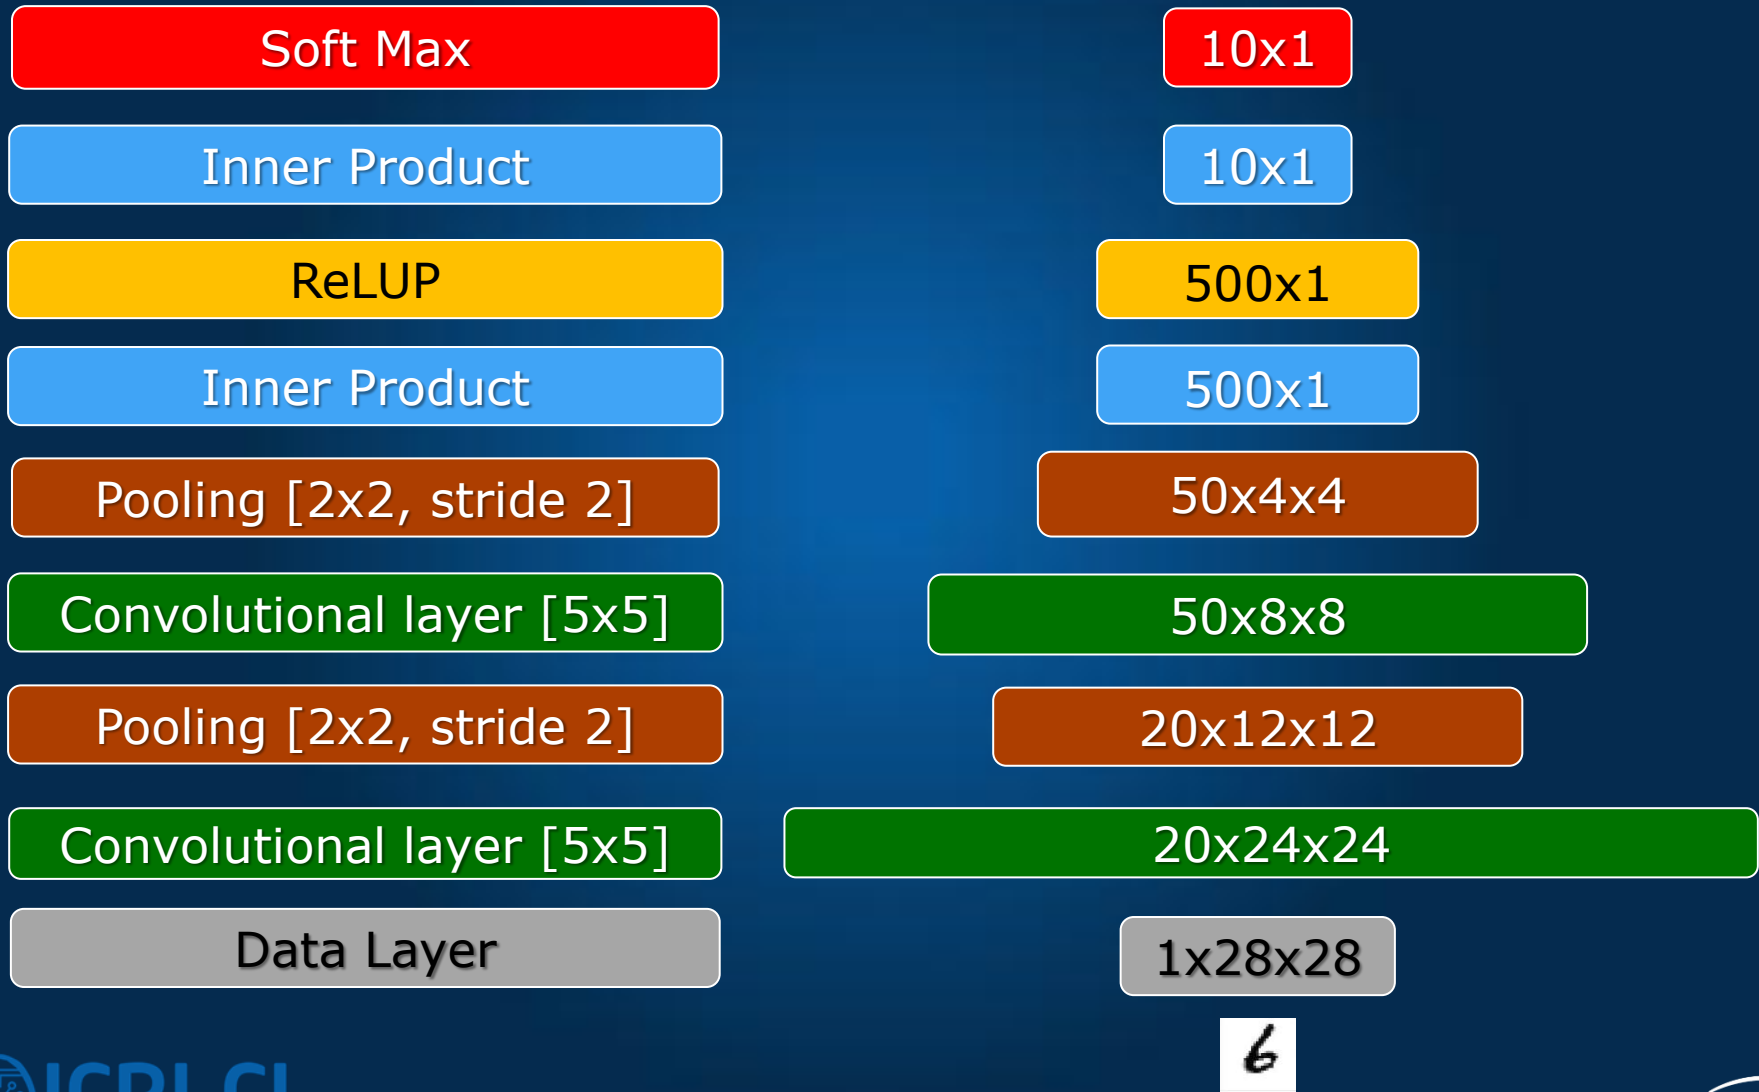
Combines softmax:

$$Y_L[i] = \exp(Y_{L-1}[i]) / (\sum(Y_L[i]));$$

with log-loss :

$$E = -\log(Y_L(\text{label}(n)))$$

LeNet topology



SOME IMPLEMENTATION DETAILS

Data Layer

All data is stored as *BLOBs* – *Binary (Basic) Large Objects*

```
class Blob {
```

```
    Blob( int num, int channels, int height, int width);
```

```
    const Dtype* cpu_data() const;
```

```
    const Dtype* gpu_data() const;
```

```
    ...
```

```
protected:
```

```
    shared_ptr<SyncedMemory> data_; // container for cpu_ / gpu_memory
```

```
    shared_ptr<SyncedMemory> diff_; // gradient
```

```
    int num_;
```

```
    int channels_;
```


```
    int height_;
```

```
    int width_;
```

```
    int count_;
```

Convolutional Layer : im2col

Conv Layer implementation is based on reduction to matrix – matrix multiply (See Chellapilla et al , “High Performance Convolutional Neural Networks for Document Processing”)



Convolutional Layer: im2col

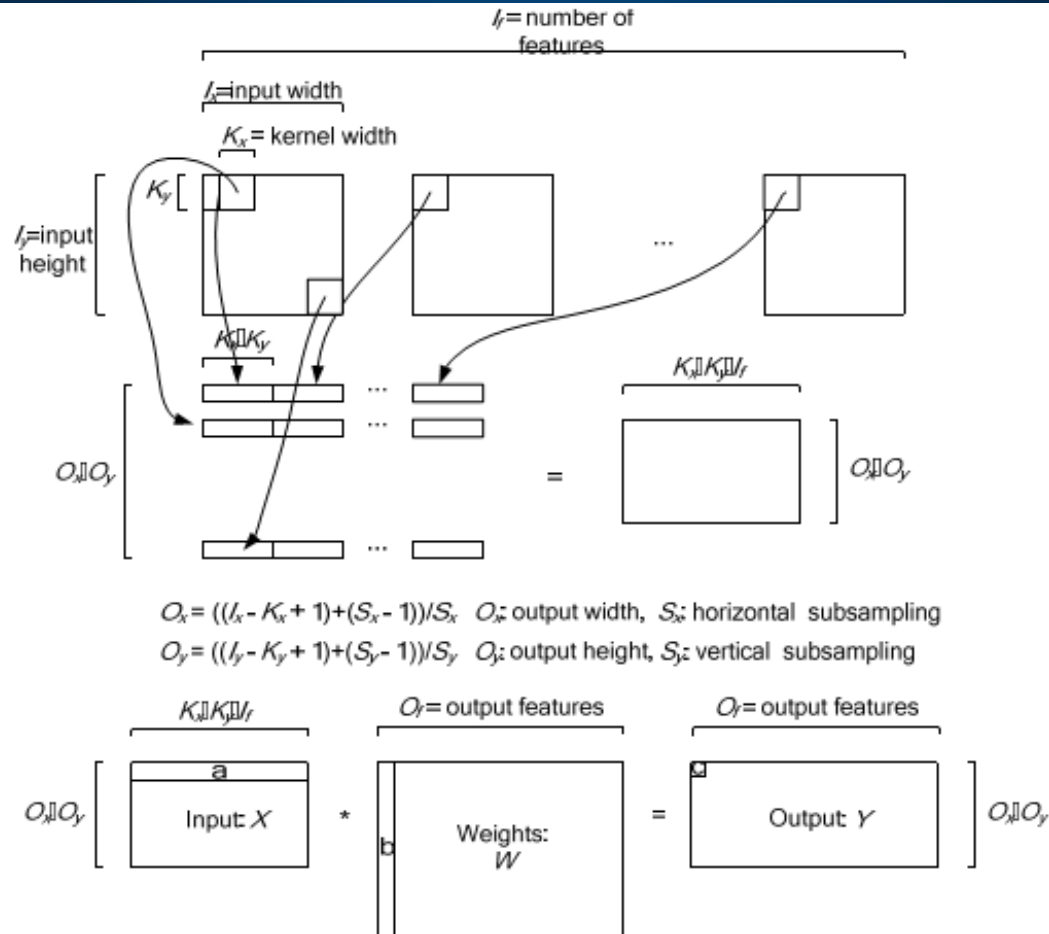
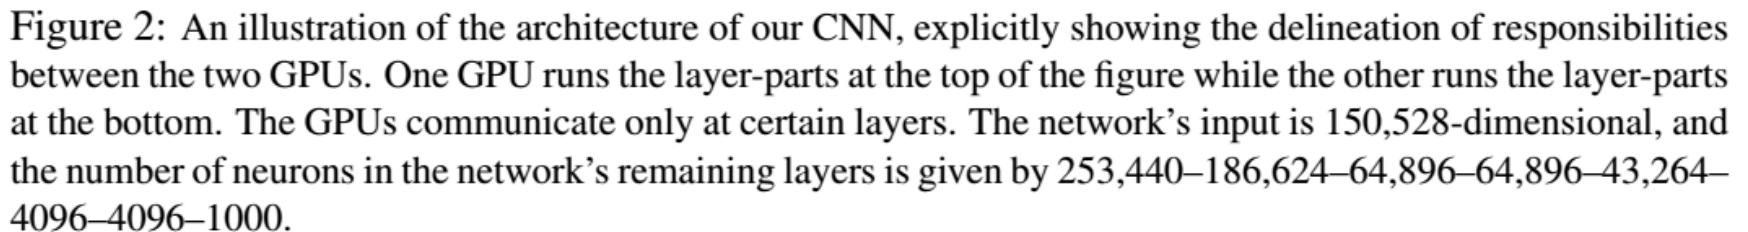


Figure 3. Unrolling the convolution operations in a convolutional layer (biases, sub-sampling, and non-linearity omitted), to produce a matrix-matrix product version.

AlexNet topology (Imagenet)



Exercises

1. Play with Mnist topologies
 - How does the net accuracy depend on topology?
2. Port one of datasets <http://deeplearning.net/datasets> :
 - NORB, SVHN, ...
3. Look at the definition of following layers:
 - sigmoid, tanh, ...